

Introduction References

- (1) Timotheou, S., Miliou, O., Dimitriadis, Y., Villagrà Sobrino, S., Giannoutsou, N., Cachia, R., Martínez Monés, A., & Ioannou, A. (2023). *Impacts of digital technologies on education and factors influencing schools' digital capacity and transformation: A literature review*. *Education and Information Technologies*, 28(6), 6695–6726.
<https://doi.org/10.1007/s10639-022-11431-8> [PMC+1](#)
- (2) Haleem, A., Javaid, M., Qadri, M. A., & Suman, R. (2022). Understanding the role of digital technologies in education: A review. *Sustainable Operations and Computers*, 3, 275–285. <https://doi.org/10.1016/j.susoc.2022.05.004>
- (3) Binkley, C. (2023, April 29). *College students are still struggling with basic math. Professors blame the pandemic*. Associated Press.
<https://apnews.com/article/college-math-test-help-6cca6a5e873d5aeb5e75b4f94125d48c>
- (4) Roumeliotis, K. I., & Tselikas, N. D. (2023). ChatGPT and Open-AI models: A preliminary review. *Future Internet*, 15(6), 192. <https://doi.org/10.3390/fi15060192>
- (5) National Assessment of Educational Progress. (2022). *NAEP mathematics: National average scores, 2022*. National Center for Education Statistics.
<https://nces.ed.gov/nationsreportcard/>
- (6) Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y. J., Madotto, A., & Fung, P. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12), 1–38. <https://doi.org/10.1145/3571730>
- (7) Oche, A. J., Folashade, A. G., Ghosal, T., & Biswas, A. (2025). A Systematic Review of Key Retrieval-Augmented Generation (RAG) Systems: Progress, Gaps, and Future Directions. *arXiv preprint arXiv:2507.18910*.
- (8) Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W.-t., Rocktäschel, T., Riedel, S., & Kiela, D. (2020, May 22). *Retrieval-Augmented Generation for knowledge-intensive NLP tasks* (arXiv Preprint No. 2005.11401). arXiv.
<https://doi.org/10.48550/arXiv.2005.1140>
- (9) Swacha, J., & Gracel, M. (2025). *Retrieval-augmented generation (RAG) chatbots for education: A survey of applications*. *Applied Sciences*, 15(8), Article 4234.
<https://doi.org/10.3390/app15084234>

Rag with LLM's

Ayala, O. M., & Béchar, P. (2024). Reducing hallucination in structured outputs via Retrieval-Augmented Generation. *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language*

Technologies (Volume 6: Industry Track), 228–238.

<https://doi.org/10.18653/v1/2024.naacl-industry.19>

Kirchenbauer, J., & Barns, C. (2024). Hallucination reduction in large language models with retrieval-augmented generation using Wikipedia knowledge. *Society*.

<https://doi.org/10.31219/osf.io/pv7r5>

Gupta, S., Ranjan, R., & Singh, S. N. (2024). A comprehensive survey of retrieval-augmented generation (RAG): Evolution, current landscape, and future directions. *arXiv*. <https://doi.org/10.48550/arXiv.2410.12837>

Gao, Y., Xiong, Y., Gao, X., Jia, K., Pan, J., Bi, Y., Dai, Y., Sun, J., Wang, H., & Wang, H. (2023). Retrieval-augmented generation for large language models: A survey. *arXiv preprint arXiv:2312.10997*, 2(1). <https://doi.org/10.48550/arXiv.2312.10997>

Shuster, K., Poff, S., Chen, M., Kiela, D., & Weston, J. (2021). Retrieval augmentation reduces hallucination in conversation. *Findings of the Association for Computational Linguistics: EMNLP 2021*, 3784–3803. <https://doi.org/10.18653/v1/2021.findings-emnlp.320>

An Introduction to RAG vs Fine-Tuning LLMs

(6) <https://aisera.com/blog/llm-fine-tuning-vs-rag/>

Retrieval-Augmented Generation: Using your Data with LLMs

<https://www.infracloud.io/blogs/retrieval-augmented-generation-using-data-with-llms/>

React and Fast API as an app development stack

React. (n.d.). *React*. <https://react.dev/>

FastAPI. (n.d.). *FastAPI*. <https://fastapi.tiangolo.com/>

Hallucinations in LLM's

(!)Holtzman, A., Buys, J., Du, L., Forbes, M., & Choi, Y. (2020). The curious case of neural text degeneration. *Proceedings of the International Conference on Learning Representations (ICLR)*. <https://arxiv.org/abs/1904.09751>

(2)Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., Chen, Q., Peng, W., Feng, X., Qin, B., & Liu, T. (2024). *A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions*. *ACM Transactions on Information Systems*, 1(1), Article 1. <https://doi.org/10.1145/3703155>

(3)Zhai, C., Wibowo, S., & Li, L. D. (2024). The effects of over-reliance on AI dialogue systems on students' cognitive abilities: A systematic review. *Smart Learning Environments*, 11(28). <https://doi.org/10.1186/s40561-024-00316-7>

Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y., Chen, D., Dai, W., Chan, H. S., Madotto, A., & Fung, P. (2024). *Survey of hallucination in natural language generation*. *ACM Computing Surveys*, 1(1), Article 1. <https://doi.org/10.1145/3571730>

Kalai, A. T., Nachum, O., Vempala, S. S., & Zhang, E. (2025, September 4). *Why language models hallucinate*. OpenAI. <https://openai.com/index/why-language-models-hallucinate/>

Hallucinations

(2) Maynez, J., Narayan, S., Bohnet, B., & McDonald, R. (2020). On faithfulness and factuality in abstractive summarization. *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, 1906–1919. <https://doi.org/10.18653/v1/2020.acl-main.173>

Zhang

Kalai, A. T., Nachum, O., Vempala, S. S., & Zhang, E. (2025, September 4). *Why language models hallucinate*. OpenAI. <https://openai.com/index/why-language-models-hallucinate/>

(3) Dziri, N., Kamaloo, E., Milton, S., Zaiane, O., Yu, M., Ponti, E. M., & Reddy, S. (2022). FaithDial: A faithful benchmark for information-seeking dialogue. *Transactions of the Association for Computational Linguistics*, 10, 1473–1490. <https://aclanthology.org/2022.tacl-1.84/>

(4) Zhai, C., Wibowo, S., & Li, L. D. (2024). The effects of over-reliance on AI dialogue systems on students' cognitive abilities: A systematic review. *Smart Learning Environments*, 11(1), 28. <https://doi.org/10.1186/s40561-024-00316-7>

(5) Khalil, M., & Er, E. (2023). Will ChatGPT get you caught? Rethinking of plagiarism detection. In M. Khalil & E. Er (Eds.), *Proceedings of the 2023 International Conference on Artificial Intelligence and Education* (pp. 1–10). Springer. https://doi.org/10.1007/978-3-031-34411-4_32

(6) Gao, C. A., Howard, F. M., Markov, N. S., Dyer, E. C., Ramesh, S., Luo, Y., & Pearson, A. T. (2023). Comparing scientific abstracts generated by ChatGPT to real abstracts with detectors and blinded human reviewers. *npj Digital Medicine*, 6(1), 75. <https://doi.org/10.1038/s41746-023-00819-6>

(7) Watts, F. M., Dood, A. J., Shultz, G. V., & Rodriguez, J.-M. G. (2023). Comparing student and generative artificial intelligence chatbot responses to organic chemistry writing-to-learn assignments. *Journal of Chemical Education*, 100(10), 3806–3817.
<https://doi.org/10.1021/acs.jchemed.3c00664>

(8) Achiam, J., Adler, S., Agarwal, S., Ahmad, L., Akkaya, I., Aleman, F. L., Almeida, D., Altenschmidt, J., Altman, S., & Anadkat, S. (2023). *GPT-4 technical report*. arXiv.
<https://doi.org/10.48550/arXiv.2303.08774>

(9) Zhang, P., Xiao, S., Liu, Z., Dou, Z., & Nie, J.-Y. (2023). Retrieve anything to augment large language models. *arXiv Preprint*. <https://doi.org/10.48550/arXiv.2310.07554>

(10) Arguson, A. C., Paculan, R. S., & Mabborang, M. P. (2024). Investigating the acceptability of students on generative AI tools: A correlation analysis among selected tertiary schools in Metro Manila, Philippines. *Journal of Artificial Intelligence, Machine Learning and Neural Network*, 4(5), 7–18.

(11) Tabuena, A. K. G. (2025). Practices of instructors in integrating artificial intelligence (AI) in assessing student performance among the community colleges in Albay. *International Journal of Research and Scientific Innovation*, 12(4), 1228–1234.
<https://doi.org/10.51244/IJRSI.2025.12040101>

(12) Garcia, M. L. (2024). Transformative impact: The evolution of learning through AI in education. *International Journal of Research and Scientific Innovation*, 11(11), 45–52.
<https://doi.org/10.51244/IJRSI.2024.111105>

(13) Espeño, S. R., Babiano, E. H., Bucoy, M. L. R., Busime, E. L., & De Borja, J. M. A. (2024). Issues and challenges of implementing special education (SPED) curriculum in the Philippines: A systematic literature review. *Jurnal Inovasi Pendidikan*, 2(4), 196–205.
<https://edukhasi.org/index.php/jip>

Rag in mitigation of Hallucinations

(1) Béchard, P., & Marquez Ayala, O. (2024). *Reducing hallucination in structured outputs via Retrieval-Augmented Generation*. arXiv.
<https://doi.org/10.48550/arXiv.2404.08189>

(2) Shuster, K., Poff, S., Chen, M., Kiela, D., & Weston, J. (2021). *Retrieval augmentation reduces hallucination in conversation*. arXiv.
<https://doi.org/10.48550/arXiv.2104.07567>

(3)Zhang, W., & Zhang, J. (2025). Hallucination mitigation for retrieval-augmented large language models: A review. *Mathematics*, 13(5), 856.

<https://doi.org/10.3390/math13050856>

(4) Wang, S.; Yu, X.; Wang, M.; Chen, W.; Zhu, Y.; Dou, Z. Richrag: Crafting rich responses for multi-faceted queries in retrieval-augmented generation. *arXiv* **2024**, arXiv:2406.12566.

(5) Song, J., Wang, X., Zhu, J., Wu, Y., Cheng, X., Zhong, R., & Niu, C. (2024). *RAG-HAT: A Hallucination-Aware Tuning Pipeline for LLM in Retrieval-Augmented Generation*. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: Industry Track* (pp. 1548-1558). Association for Computational Linguistics.

(6)Sun, Z., Zang, X., Zheng, K., Song, Y., Xu, J., Zhang, X., Yu, W., & Li, H. (2024). *ReDeEP: Detecting hallucination in retrieval-augmented generation via mechanistic interpretability*. *arXiv*.

<https://doi.org/10.48550/arXiv.2410.11414>

(7)Hu, H., He, C., Xie, X., & Zhang, Q. (2024). *LRP4RAG: Detecting Hallucinations in Retrieval-Augmented Generation via Layer-wise Relevance Propagation*. *arXiv*.

<https://doi.org/10.48550/arXiv.2408.15533> arxiv.org

Methodology

Semantic Chunking

Kamradt, G. (n.d.). *5 levels of text splitting* [Jupyter notebook]. Full Stack Retrieval / GitHub. Retrieved September 15, 2025, from

https://github.com/FullStackRetrieval-com/RetrievalTutorials/blob/main/tutorials/LevelsOfTextSplitting/5_Levels_Of_Text_Splitting.ipynb

LangChain. (n.d.). *Semantic chunker*. LangChain Documentation. Retrieved September 15, 2025, from https://python.langchain.com/docs/how_to/semantic-chunker/